

Coding

Terminal-Bench 4.0

- Description: The 4.0 release of Terminal-Bench, developed by Stanford University researchers, the Laude Institute, and the open source community. Covers software engineering, system administration, data processing, model training, and security, with each task graded by its own verification suite.
- Leaderboard: <https://www.tbench.ai/?version=4>
- Dataset: <https://github.com/harbor-framework/terminal-bench>
- Implementation:
 - We evaluate the full Terminal-Bench 4.0 dataset (66 tasks) using the mini-swe-agent harness, with pass@1 scoring averaged over 3 repeats per task
 - Each task has its own set of tests. We follow the Terminal-Bench methodology: a task passes only if every test passes, grading runs in each task's own separate verifier container isolated from the agent's environment, and a verifier that exceeds its timeout counts as a failure
 - We apply the following constraints on evaluations for the agent:
 - Maximum agent steps are limited to 500
 - Task timeouts and sandbox resources follow the upstream task definitions
 - All other agent configuration follows mini-swe-agent defaults, including the interactive mini config and prompts, the native bash tool, and no context compaction or summarization: the agent always sees its full transcript

SciCode

- Description: Python programming to solve scientific computing tasks.
- Paper: <https://arxiv.org/abs/2407.13168>
- Dataset: <https://scicode-bench.github.io/>
- Implementation:
 - We test with scientist-annotated background information included in the prompt
 - We report sub-problem level scoring
 - Pass@1 evaluation criteria
 - SciCode step scripts are graded on isolated executors with a 300-second execution timeout (dataset v1.0.1)

HARNESS: <https://github.com/SWE-agent/mini-swe-agent.git>